

Analytical Data Modeling (With On-Chain Data)

@sha2fiddy

Engineering Manager, Data Analytics

Foundry Digital

btc++ Austin 2025

Analytics Engineering

What: transform raw data into clean, reliable data models

Analogy: raw tables are ingredients; models are recipes


Why: speed of analysis, consistency, & reusability

How: applying software engineering principals and technologies to data modeling & analysis

Data Engineers: build data ingestion pipelines & infrastructure

Data Analysts: create queries, reports, & dashboards

Analytics Engineers: build data transformation pipelines with software engineering best practices, but primarily to enable analysts



Data Modeling Concepts

Data Granularity: level of detail of a data model

- *week/day/hour*
- *block/address/transaction*
- *pool/subaccount/worker*

Natural vs Surrogate Key: a unique identifier from the data source vs one generated within the data warehouse

Change Data Capture (CDC): systematically record how & when categorical data changes over time (e.g. name change)

Related: Slowly Changing Dimensions (SCDs)

Normalization vs Denormalization: separating categorical and numerical data vs combining them into flat dataset models



Star Schema Models

Star Schema: a central fact table which joins to one or more dimension tables (**normalized**)

Dimension (Dim): categorical data that is used to aggregate, filter, and sort by (one row per dim entity)

- *date, pool, customer, miner model*

Fact: numerical measurements or events to be aggregated (one row per level of detail)

- *revenue, hashrate, price, worker count*

Pros: reduced storage, reduced code duplication, central sources of truth

Cons: joins required, less intuitive for business users



OBT Models

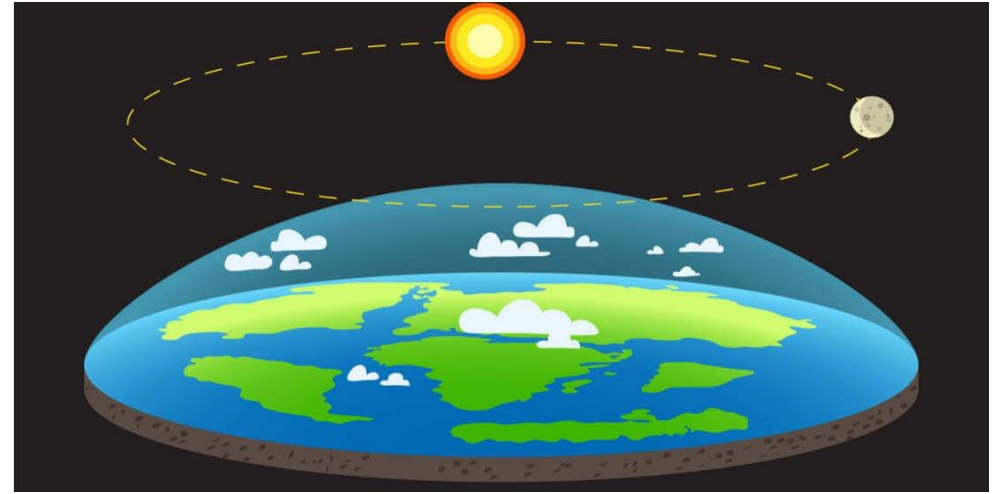
OBT: Operational BI Table (aka One Big Table)

Flat standalone models (**denormalized**), containing all attributes that would otherwise be broken into dims & facts

Pros: fewer joins required, more intuitive for business users

Cons: increased storage, code duplication, can introduce many categorical flags or partitions needed for filters

Best of Both: build star schemas for primary dims & facts, and application-specific OBTs which read from those dims & facts



Data Warehouse Tools & Tips

Data Build Tool (dbt): data modeling framework that brings automation, testing, & jinja templating to data transformation (core product is open source)

Analogy: like a frontend application framework, but for analytics engineering

Columnar Databases: column-oriented databases are optimized for aggregating large volumes of data (vs traditional RDS systems which are transactional)

- **SaaS:** Snowflake, BigQuery, Databricks
- **Open Source:** DuckDB, ClickHouse, TimescaleDB

Extract, Load, Transform (ELT): ingest raw data to a central storage layer first, then perform all transformations (vs traditional ETL approach)

Data Validation: automated data audits & tests on raw and modeled layers help ensure data integrity & reliability



Model Examples

DIM Categorical

| dim.block |
|--------------------------|
| block_id |
| block_hash |
| blockheight |
| coinbase_address |
| coinbase_tag |
| is_stale |
| is_subsidy_halving |
| is_difficulty_adjustment |

| dim.date |
|--------------|
| date_id |
| date |
| year |
| quarter |
| month |
| day_of_week |
| day_of_month |
| week_of_year |
| is_weekend |

| dim.pool |
|-------------------|
| pool_id |
| pool_name |
| pool_url |
| start_date |
| is_antpool_friend |

| dim.coin |
|----------------|
| coin_id |
| coin_name |
| start_date |
| is_second_best |

FACT Metric

| fact.block |
|-------------------|
| block_id |
| coin_id |
| pool_id |
| date_id |
| timestamp |
| difficulty |
| block_size |
| transaction_count |
| reward_subsidy |
| reward_tx_fee_sum |

| fact.price_1d |
|---------------|
| date_id |
| coin_id |
| price_open |
| price_close |
| price_low |
| price_high |
| trade_volume |

| fact.network_stats_1d |
|-------------------------|
| date_id |
| coin_id |
| difficulty_weighted_avg |
| block_count |
| estimated_hashrate |
| reward_subsidy_sum |
| reward_tx_fee_sum |
| reward_tx_fee_avg |

| fact.pool_stats_1d |
|----------------------|
| date_id |
| pool_id |
| coin_id |
| estimated_hashrate |
| reported_hashrate |
| expected_block_count |
| block_count |
| mining_luck |
| reward_tx_fee_avg |

OBT Dataset

| obt.block |
|----------------------|
| <dim.block> |
| <dim.coin> |
| <dim.date> |
| <dim.pool> |
| <fact.bitcoin_block> |

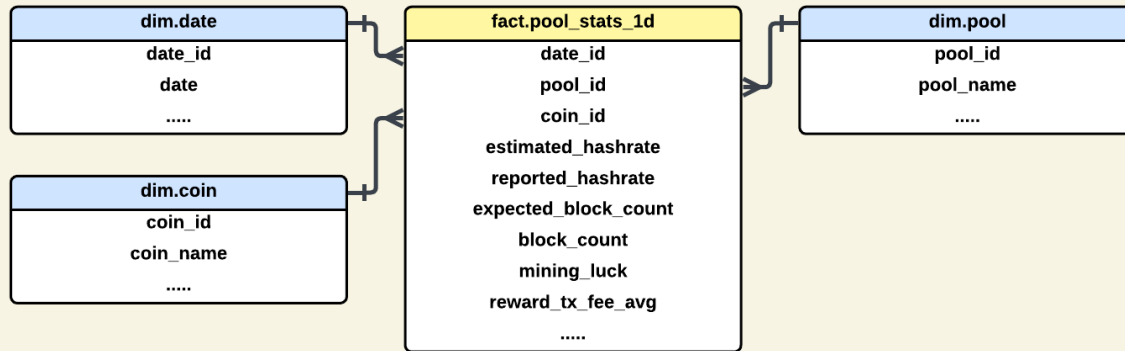
| obt.pool_stats_all_time |
|-------------------------|
| <dim.pool> |
| <dim.coin> |
| <fact.bitcoin_block> |
| <fact.network_stats_1d> |
| <fact.pool_stats_1d> |

| obt.network_stats_1d |
|-------------------------|
| <dim.date> |
| <dim.coin> |
| <fact.network_stats_1d> |
| <fact.spot_price_1d> |

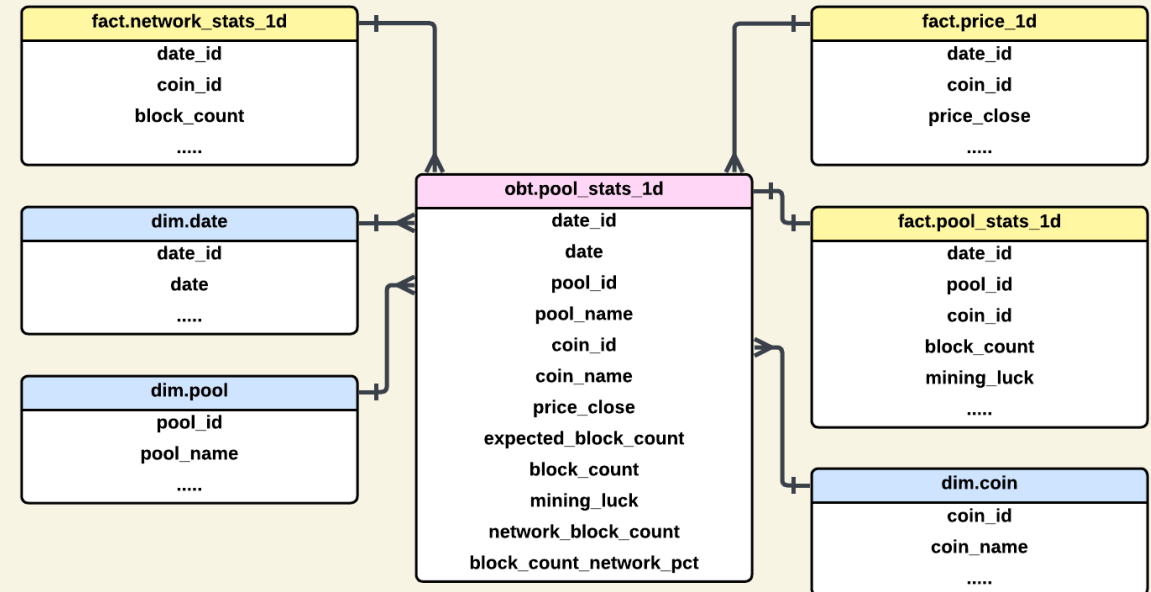
| obt.pool_stats_1d |
|-------------------------|
| <dim.date> |
| <dim.pool> |
| <dim.coin> |
| <fact.network_stats_1d> |
| <fact.pool_stats_1d> |
| <fact.spot_price_1d> |

Schema Examples

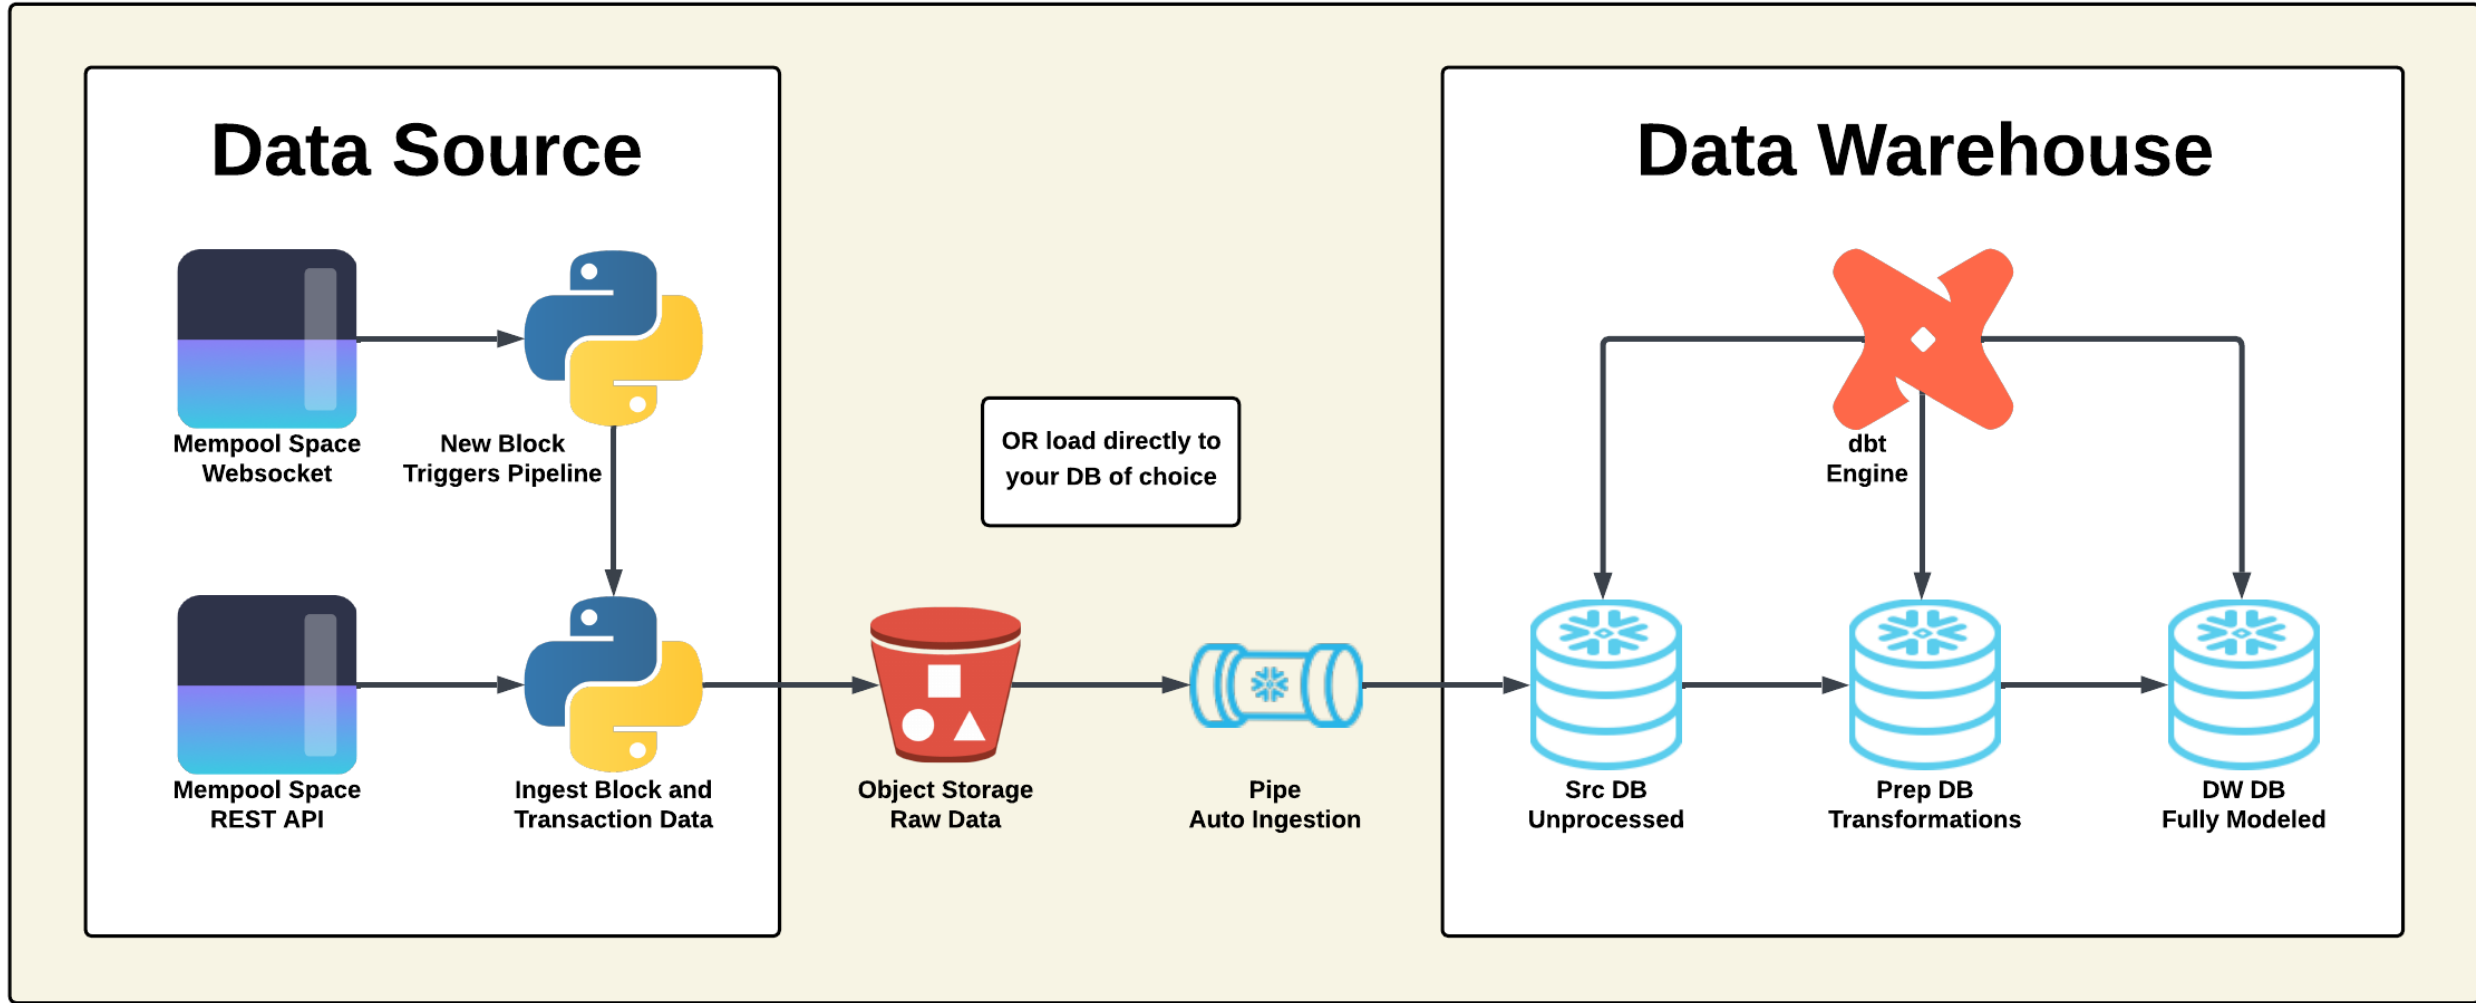
Star Schema Normalized



OBT Denormalized



Pipeline Architecture Example



Workshop

Star Schemas:

- **Dims:** date, block, pool
- **Facts:** block, network stats, pool stats, price

OBTs: block, network stats, pool stats

Metrics: estimated network hashrate, hashprice, pool mining luck

Repo: <https://github.com/sha2fiddy/btcpp-demo-2025>

